

An Autonomous Agent that Derives ADC Linker Rules from the Literature, Designs the Molecules, and Flags When the Field Disagrees

Elia Savino¹, Joshua W. Sin^{2,3}, Morgan G. L. Reigner¹, Miquèl À. Pérez-Puigdomènech², Derk H. W. ten Klooster¹

¹Noël Research Group, Van 't Hoff Institute for Molecular Sciences, University of Amsterdam, The Netherlands

²Laboratory of Artificial Chemical Intelligence (LIAC), EPFL, Lausanne, Switzerland

³Process Chemistry & Catalysis, F. Hoffmann-La Roche AG, Basel, Switzerland

Abstract

We report an autonomously derived, falsifiable prediction for antibody–oligonucleotide conjugate (ARC) design: a protease-cleavable Val-Cit linker is a viable alternative to the rigid non-cleavable sulfo-SMCC standard, because the modern clinical ARC literature has already moved that way. An agent reached this independently—it reads 30 primary papers and derives a design rule per payload class with a retrieval-grounded language model, and a leave-one-paper-out test shows the ARC rule is contested: withhold the rigid-non-cleavable siRNA paper and the agent derives protease-cleavable Val-Cit from the clinical antibody-oligonucleotide literature (DYNE-101/251, AOC-1001), lowering its confidence to 0.33 because the grounded evidence genuinely splits. Each rule parameterises a REINVENT4 objective that generates the linkers, and a Val-Cit ARC linker generated under the derived cleavable rule is recognised by cathepsin B (predicted K_d 10 nM), as tightly as the clinical Val-Cit substrate (35 nM) and unlike the rigid ARC design (109 nM)—an in-silico feasibility proof for the prediction. By contrast the cytotoxin and ISAC rules are robust (they recover at 0.97 and 0.83 when a paper is withheld), so the agent’s confidence tracks the literature’s actual consensus. Every reported value is provenance-tagged and a gate blocks placeholders from results, so the prediction is a grounded, calibrated conclusion a lookup-table system could not produce.

1 Introduction

Antibody conjugates couple antibody selectivity to a potent cargo, and the linker sets their therapeutic index [1, 3]. The linker requirement inverts with payload class: cytotoxins favour cleavable release [2], ISACs demand plasma stability [6, 7], and antibody–oligonucleotide conjugates (ARCs) were reported to favour rigid non-cleavable linkers [5]. An autonomous scientist should derive such rules from the literature, use them to design molecules, and report how confident it is—including recognising where the literature genuinely conflicts. We show all three, and the last surfaces a real, unresolved tension in ARC linker design (Figure 1) that we sharpen into a testable prediction:

Hypothesis (agent-derived, falsifiable): for antibody–oligonucleotide conjugates, a protease-cleavable Val-Cit linker matches or outperforms the rigid non-cleavable sulfo-SMCC standard on protease-mediated payload release. We derive it from the literature and test its structural feasibility in silico.

2 Methods

2.1 Reasoning that drives generation

For each payload class an agent (Claude Opus, always on) queries a store of 30 primary-literature documents (the full corpus is listed in SI), retrieves passages with provenance, and extracts conjugate exemplars that must each cite a retrieved passage. It derives a rule (cleavage, rigidity, plasma-stability) from the grounded exemplars—with no hardcoded fallback, so a failed call is an explicit abstain—and the rule compiles into an ADC goal profile (cleavage regime, rotatable-bond window, per-term weights) that parameterises the REINVENT4 LinkInvent objective. That objective then generates the linkers (staged learning, remote GPU); the class trigger warhead (Val-Cit-PABC, Val-Ala-PABC or the rigid sulfo-SMCC cap) is welded into every molecule (Figure 3).

2.2 The disagreement-aware confidence

Confidence in a derived rule is defined explicitly as evidence strength times consensus. Let n be the number of grounded exemplars, and among those whose linker is cleavage-labelled let n_c favour cleavable and n_n non-cleavable

One molecule, end to end: a retrieved quote grounds the rule that generates the linker that cathepsin B recognises




Figure 1: One molecule, end to end: a retrieved quote grounds the cytotoxin rule that generates a Val-Cit linker that cathepsin B recognises. Handle (blue), scissile bond (red) and spacer (green) are SMARTS-detected.

release. With consensus $\kappa = \max(n_c, n_n)/(n_c + n_n)$,

$$C = \min(0.97, E \cdot P), \quad E = \min(1, \frac{n}{6}), \quad P = \max(0, 2\kappa - 1). \quad (1)$$

E rewards evidence volume (saturating at six grounded exemplars); P rewards consensus, falling linearly from 1 at unanimity to 0 at an even split, so a contested field drives C down regardless of how many papers were read. A rule is flagged *contested* when $n_c + n_n \geq 3$ and $\kappa < 0.7$. This makes every confidence in the paper derivable: e.g. the full-corpus cytotoxin evidence is 12:3 cleavable:non-cleavable ($\kappa = 0.80$, $E = 1$) so $C = 0.60$, whereas withholding the review that supplies those non-cleavable examples leaves 8:0 ($\kappa = 1$) so $C = 0.97$ —the rise reflects a genuinely more consistent evidence subset, not a change in how many papers were seen.

2.3 Scorers and provenance

Synthesizability is the continuous Ertl SA_Score (the metric the generator optimises); the step-count survives only as a reported complexity index, not a route. Plasma stability is mechanism-resolved. Every value that reaches the manuscript is tagged *measured* (a real Boltz or REINVENT run), *llm*, or *heuristic*; a gate blocks any placeholder from a results claim (SI provenance table).

2.4 Held-out test and co-fold

We withhold a class’s key paper, re-derive the rule from the remaining corpus, record the prediction before revealing the withheld paper, and report agree/disagree with its evidence composition. We co-fold each designed linker with its real payload against cathepsin B (Boltz-2) as two co-present ligands—a recognition/plausibility check, not a covalent conjugate and not a claim of cleavage. Boltz-2 affinity predictions carry roughly an order-of-magnitude (~ 1 log-unit) uncertainty, so we read them only as regime separators, not point values.

3 Results

3.1 The rules generate, and steer, the molecules

Each derived rule compiles into a distinct objective (Figure 3, left) that generates its linkers: the designs separate cleanly by class along the two rule knobs (Figure 3, right), ARC in a rigid non-cleavable region (4–6 rotatable bonds), cytotoxin and ISAC as cleavable flexible peptides (12–16). The molecules are made by the reasoning, not labelled after the fact.

3.2 Held-out test: robust rules and a contested one

The leave-one-paper-out test discriminates (Figure 2, Table 1). The cytotoxin and ISAC rules are robust: their grounded exemplars are unanimously cleavable, so withholding a paper leaves the derived rule and its high confidence intact. ISAC’s robustness, however, reflects evidence *volume*, not tested consensus: with a single cleavage-labelled exemplar ($n_c + n_n = 1$) it cannot be flagged contested by construction, so its 0.83 is a lower bound on uncertainty. That same thin evidence let the model downgrade ISAC’s plasma-stability priority to *standard* (from zero stability-labelled exemplars), which would have shipped a TLR7/8 agonist—whose premature systemic release drives cytokine-release syndrome—at the lowest stability weight of the three classes; a low-evidence safety gate floors it to the encoded domain prior instead (paramount, Figure 3; SI), extending the confidence machinery to the safety-critical stability




Figure 2: Held-out test. Top: full-corpus vs held-out confidence (disagreement-aware); the ARC rule flips (red) while cytotoxin/ISAC recover (green). Bottom: the grounded-evidence split that explains it—ARC is 4:2 (contested), the others unanimous.

4 Discussion

The agent reads the literature, derives rules that generate synthesisable molecules, and reports calibrated confidence that falls when the evidence divides—turning “contested” from a caption into a computed output. The ARC result is the payoff: a grounded, non-obvious prediction that clinical antibody-oligonucleotide conjugates favour cleavable Val-Cit, against the rigid-non-cleavable rule of the early siRNA literature, now with in-silico feasibility support. We report confidence both on the full corpus and on the held-out subset and label them as such: the full-corpus value can be the lower of the two (cytotoxin 0.60 vs held-out 0.97) where broad retrieval mixes in less-consistent exemplars that the disagreement-aware metric rightly penalises, while the held-out subset is more focused. *Limitations.* The complexity index is a count-based proxy, not a retrosynthetic route (AiZynthFinder is the planned upgrade); the co-fold places linker and payload as co-present ligands, not a covalent construct; the retrieval is lexical; and generation samples 25 designs per class (15 shortlisted, five per class). Each is tagged, not hidden.

5 Conclusion

An autonomous agent derived payload-class ADC-linker rules from 30 papers, used them to generate synthesisable linkers, and, under a held-out test, distinguished robust rules from a genuinely contested one—surfacing an unresolved question in ARC linker design and lowering its confidence where the literature divides. The approach points toward an autonomous medicinal-chemistry scientist that designs and knows what it does not know.

Acknowledgements

AI tools (Anthropic Claude) were used for retrieval-grounded literature reasoning, code, and manuscript drafting under author supervision. Per ICMJE and publisher policy, AI tools are not listed as authors; all scientific decisions and the final text are the authors’ responsibility.

References

- [1] Su, Z.; et al. Antibody–drug conjugates: Recent advances in linker chemistry. *Acta Pharm. Sin. B* **2021**, 11 (12), 3889–3907.




Figure 3: The rules steer the chemistry. Left: the compiled objective weights—a deterministic rule-compilation (a heuristic keyed by the agent’s derived category, effectively three buckets, not continuously tuned). The “rigidity” term rewards *fewer* rotatable bonds; the real separator is the per-class rotatable-bond setpoint ($\text{rot} \leq 5/10/14$), shown by each class. ISAC’s stability weight is floored to paramount by the low-evidence safety gate. Right: the generated designs separate by class, driven mainly by the rigidity setpoint (x); y is the per-molecule Ertl SA subscore (a real quantity, not jitter).




Figure 4: Stress-test of the confidence metric: varying retrieval depth (top- k) leaves each class’s confidence stable—the ARC rule stays contested (≈ 0.33) at every k . Reproducible, not a lucky k .




Figure 5: Cathepsin-B co-fold (affinity indexes recognition, not cleavage). The cleavable Val-Cit designs (cytotoxin, and the held-out-derived ARC linker) are recognised in the clinical-substrate regime; the rigid ARC design binds $>10\times$ weaker—the in-silico feasibility proof.

- [2] Balamkundu, S.; Liu, C.-F. Lysosomal-Cleavable Peptide Linkers in Antibody–Drug Conjugates. *Biomedicines* **2023**, 11 (11), 3080.
- [3] Peng, B.; et al. Stimulus-cleavable chemistry in controlled drug delivery. *Chem. Soc. Rev.* **2021**, 50, 4737–4762.
- [4] Su, D.; Zhang, D. Linker Design Impacts ADC Pharmacokinetics and Efficacy. *Front. Pharmacol.* **2021**, 12, 687926.
- [5] Antibody–siRNA conjugates for tumor cell gene silencing without cationic assistance. *Bioconjugate Chem.* **2025**. DOI: 10.1021/acs.bioconjchem.5c00212.
- [6] Immune-stimulating antibody conjugates elicit robust myeloid activation (TLR7/8 ISAC). *Mol. Pharmaceutics* **2022**. DOI: 10.1021/acs.molpharmaceut.2c00392.
- [7] Design of TLR7/8 agonist immune-stimulating antibody conjugates with tuned linker cleavability. *J. Med. Chem.* **2025**. DOI: 10.1021/acs.jmedchem.5c01908.
- [8] Passaro, S.; et al. Boltz-2: Towards accurate and efficient binding-affinity prediction. **2025**.